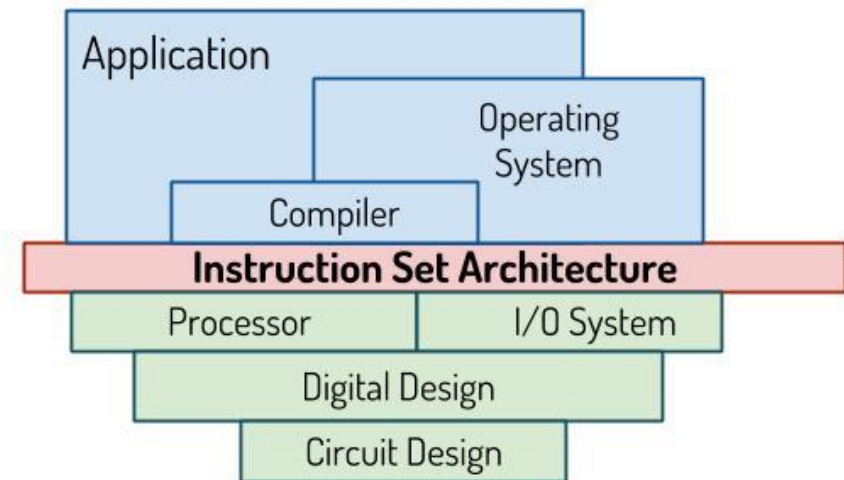

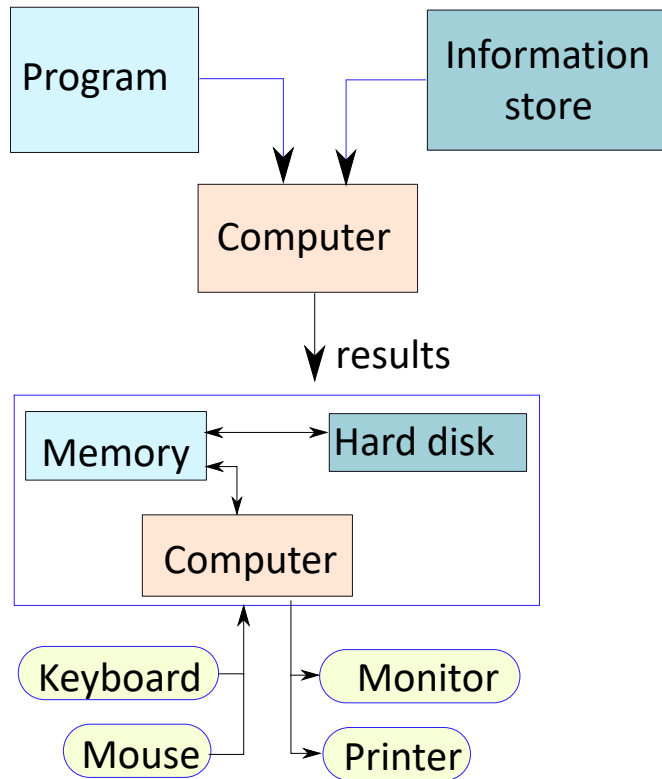


COMPUTER ARCHITECTURE

A Quick Review

The Computer

- Architecture
 - The view of a computer as presented to software designers
- Organization
 - Implementation of a computer in hardware



The Theory


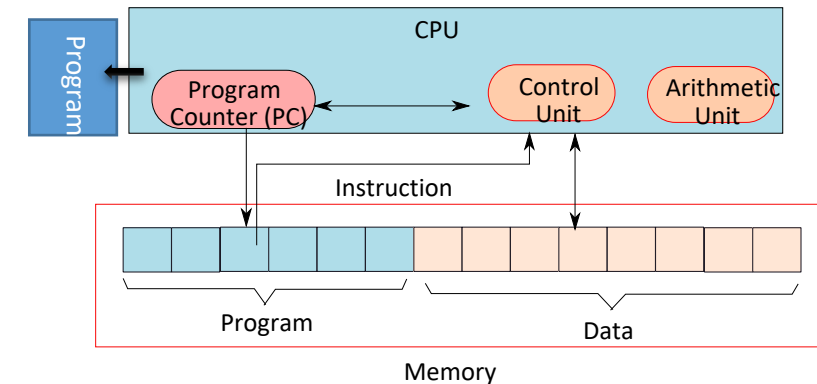
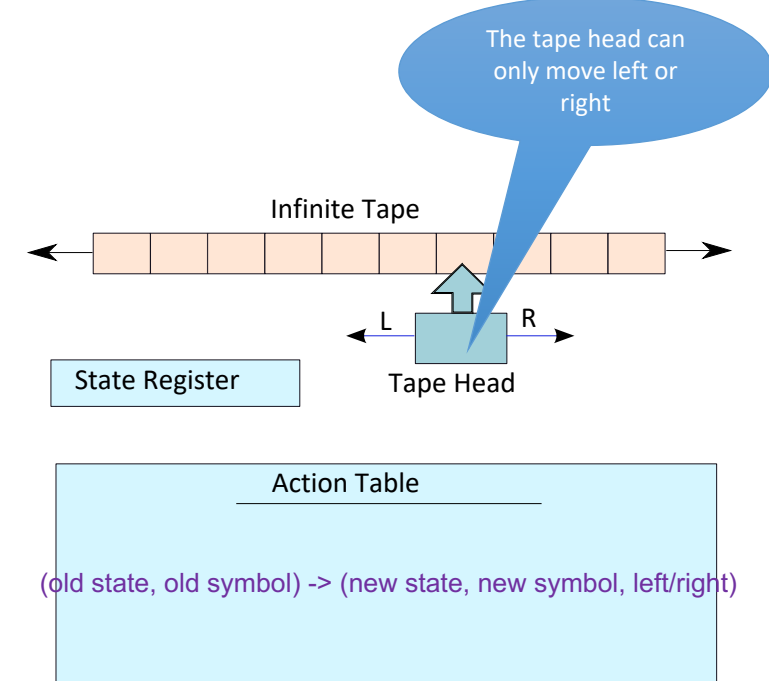

- Start at the Beginning

- Turing Machines

- * Based on the symbol in the cell, and its current state, the Turing machine computes the transition :
 - * Computes the **next state**
 - * Overwrites the symbol in the **cell** (or keeps it the same)
 - * Moves to the left or right by 1 **cell**
 - * Extremely simple, and extremely powerful

- Real Architectures

- Von Neumann
 - Harward



Instruction Set Architecture

- The Language of Instructions

- Interface between Hardware and Software

- A compiler converts a program into machine instructions in the given ISA
- The processor executes the instructions in the ISA
 - CISC/RISC

- Defined by Registers, Instructions and addressing Modes

- Machine instruction typically contains

- an opcode
- one or more source operands
- possibly a destination operand

```
addi $t0,$s0,61 # set $t0 to ($s0)+61
xori $t0,$s0,0x00ff # set $t0 to ($s0) ⊕ 0x00ff
lw $t0,40($s3) # load mem[40+($s3)] in $t0
sw $t0,A($s3) # store ($t0) in mem[A+($s3)]
j verify # go to mem loc named "verify"
```

High-level language statement:

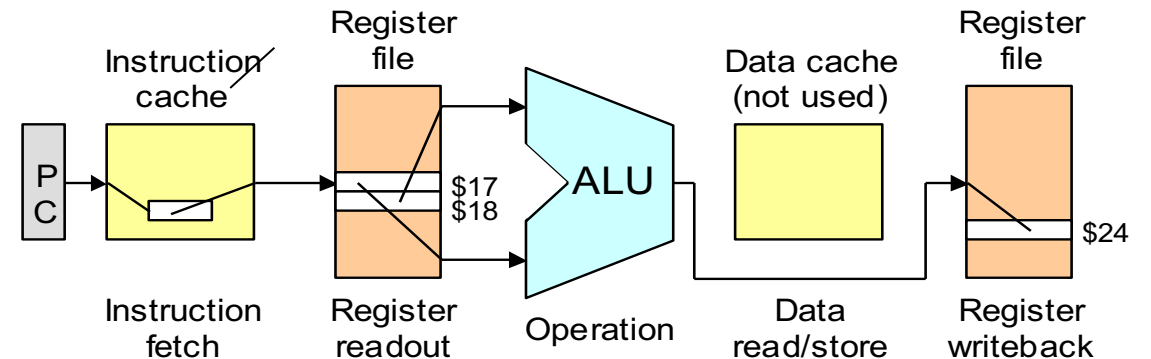
$a = b + c$

Assembly language instruction:


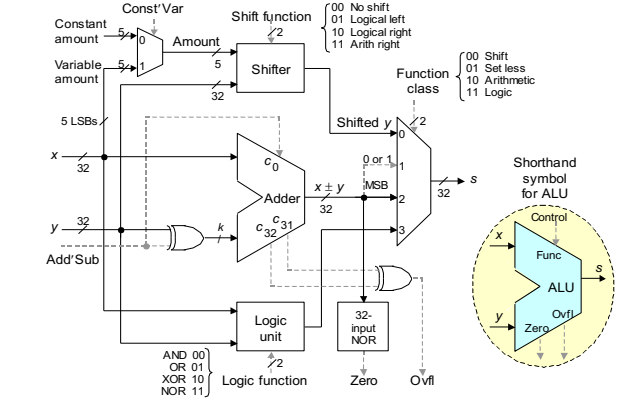
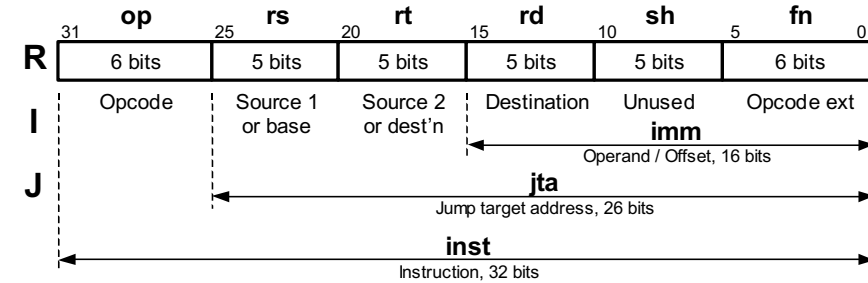
add \$t8, \$s2, \$s1

Machine language instruction:

000000 10010 10001 11000 00000 100000
ALU-type Register Register Register Unused Addition
instruction 18 17 24 opcode



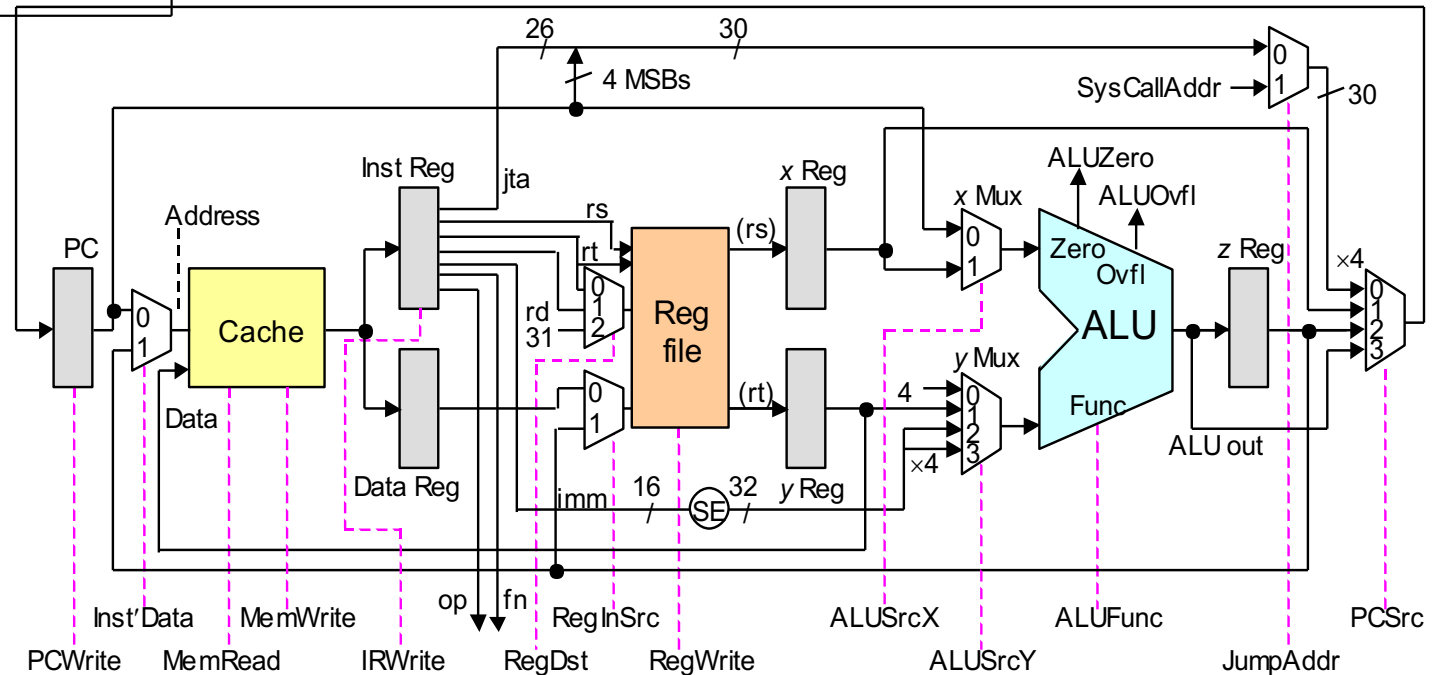
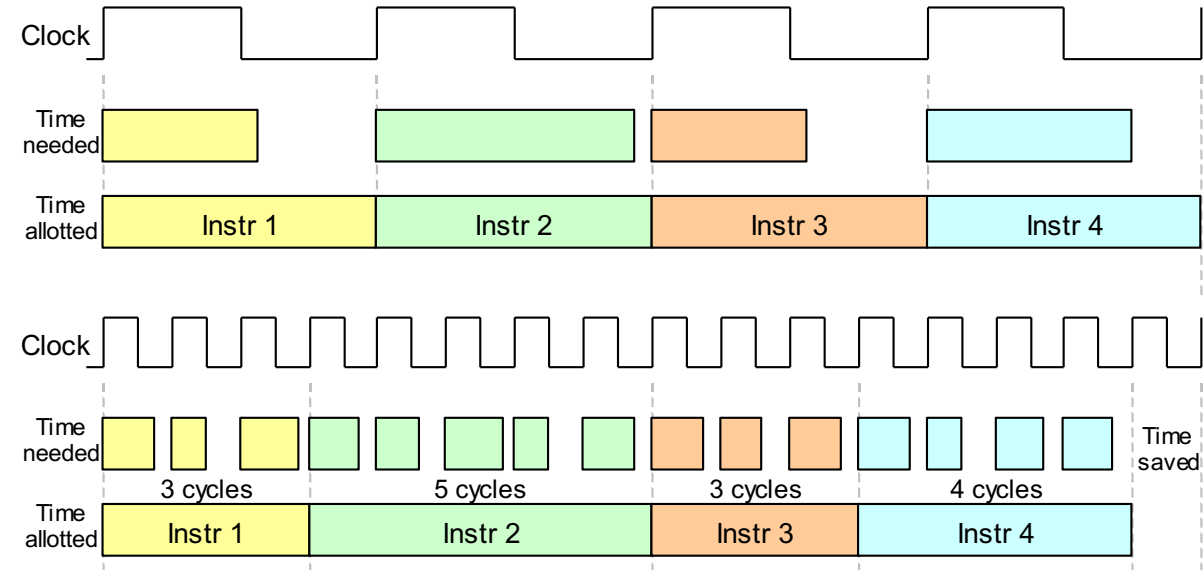
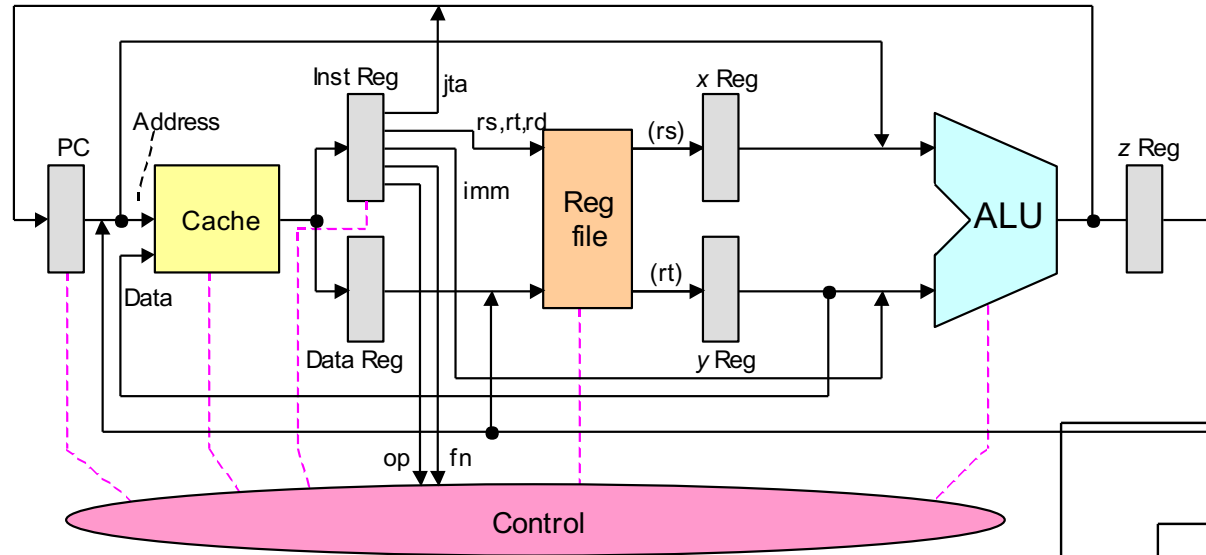
- Abstract view of the instruction execution unit



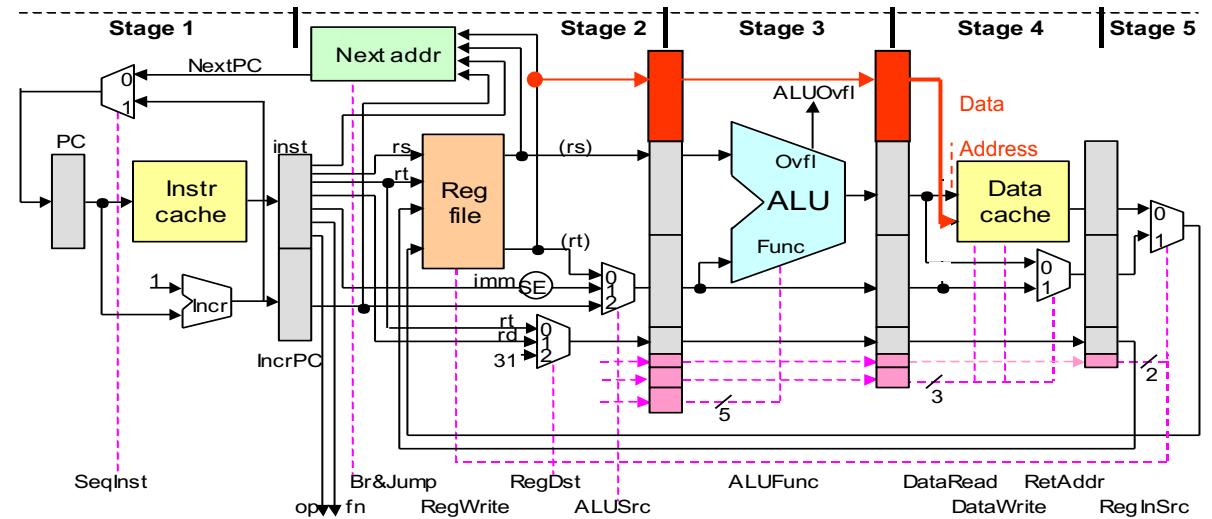
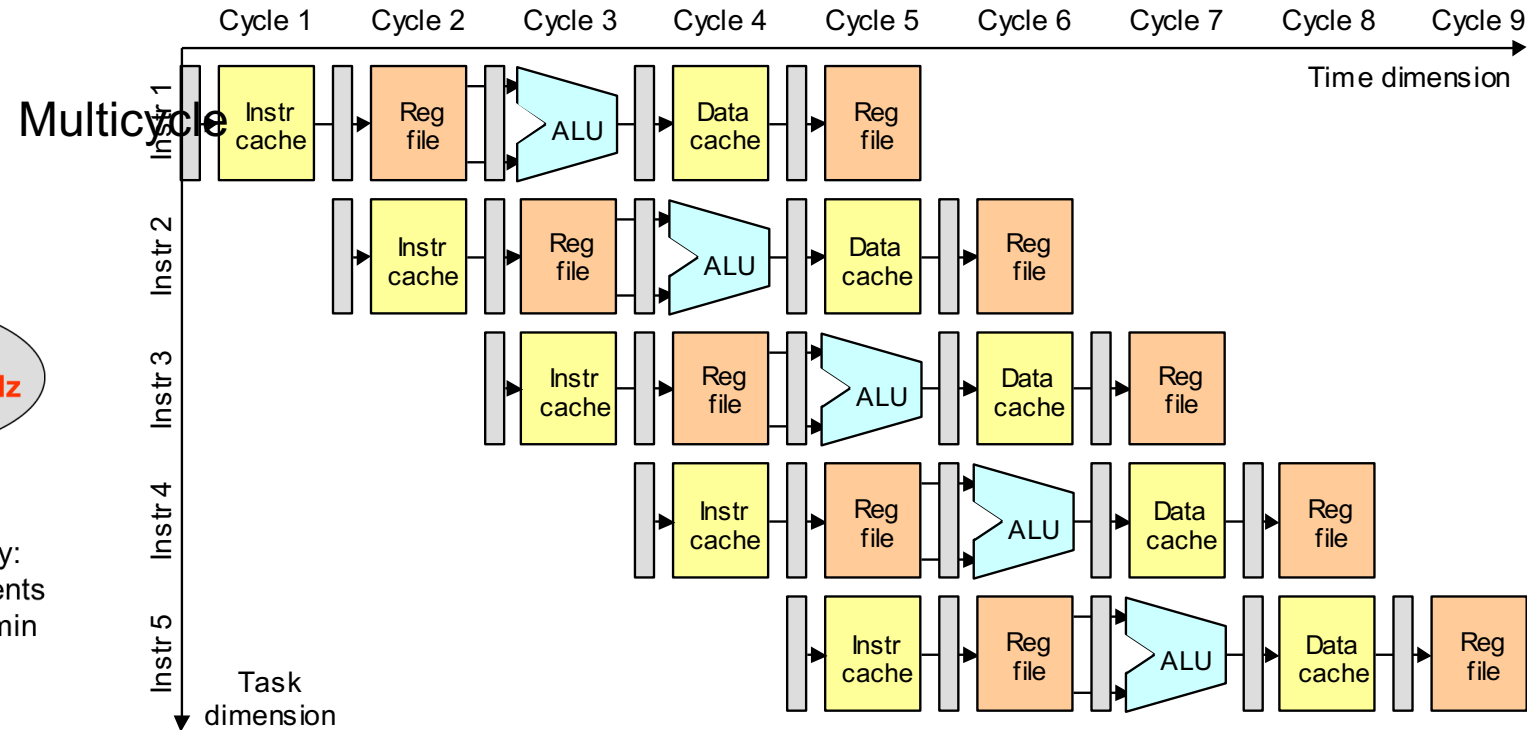
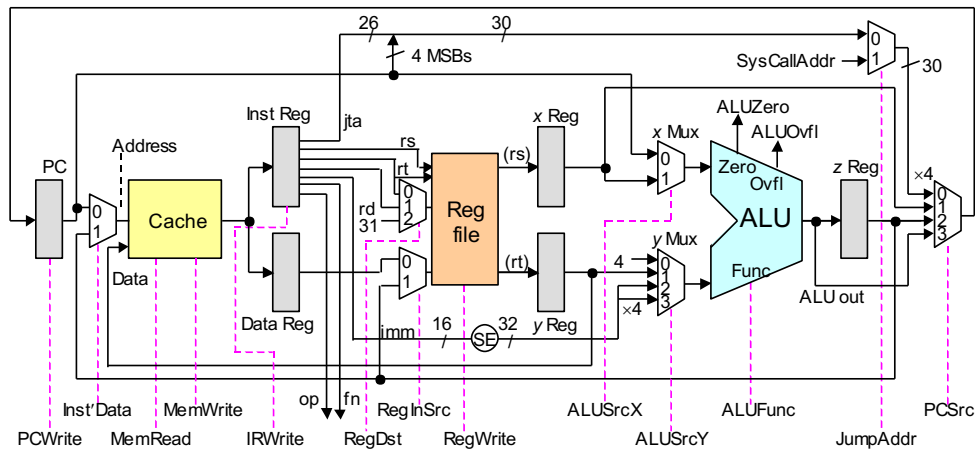
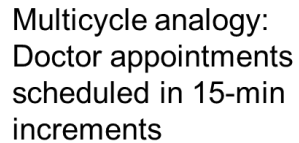
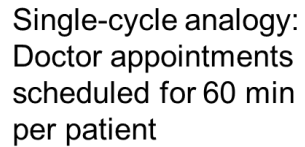
Reg write- back

Multi Cycle MicroArchitecture

- Improve Performance



- CPI + High Clock Rates



Performance Equation

• How do we Evaluate

Performance = 1 / Execution time *simplified to* 1 / CPU execution time

CPU execution time = Instructions × CPI / (Clock rate)

Performance = Clock rate / (Instructions × CPI)

Instruction access	2 ns
Register read	1 ns
ALU operation	2 ns
Data cache access	2 ns
Register write	<u>1 ns</u>
Total	8 ns

Single-cycle clock = 125 MHz

R-type	44%	6 ns
Load	24%	8 ns
Store	12%	7 ns
Branch	18%	5 ns
Jump	2%	<u>3 ns</u>

Weighted mean \cong 6.36 ns

R-type	44%	4 cycles
Load	24%	5 cycles
Store	12%	4 cycles
Branch	18%	3 cycles
Jump	2%	3 cycles

Multi-cycle clock = 500 MHz

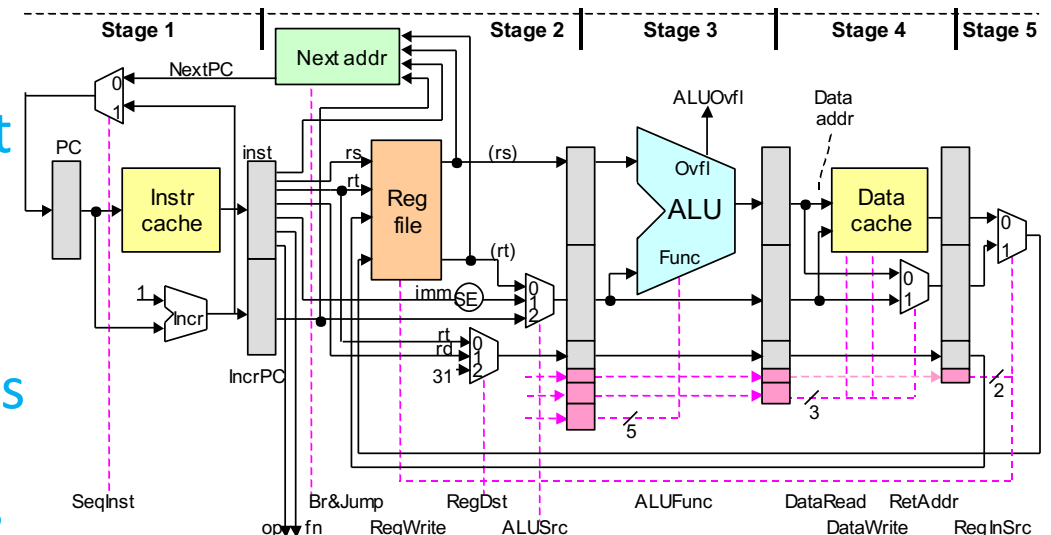
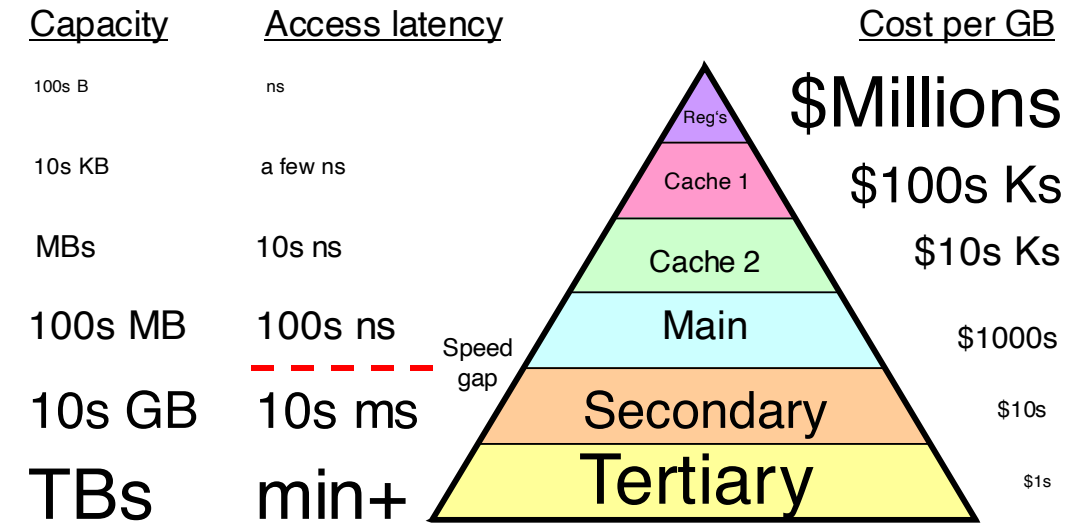
Contribution to CPI

R-type	$0.44 \times 4 = 1.76$
Load	$0.24 \times 5 = 1.20$
Store	$0.12 \times 4 = 0.48$
Branch	$0.18 \times 3 = 0.54$
Jump	$0.02 \times 3 = 0.06$

Average CPI \cong 4.04

Need for a Memory Hierarchy

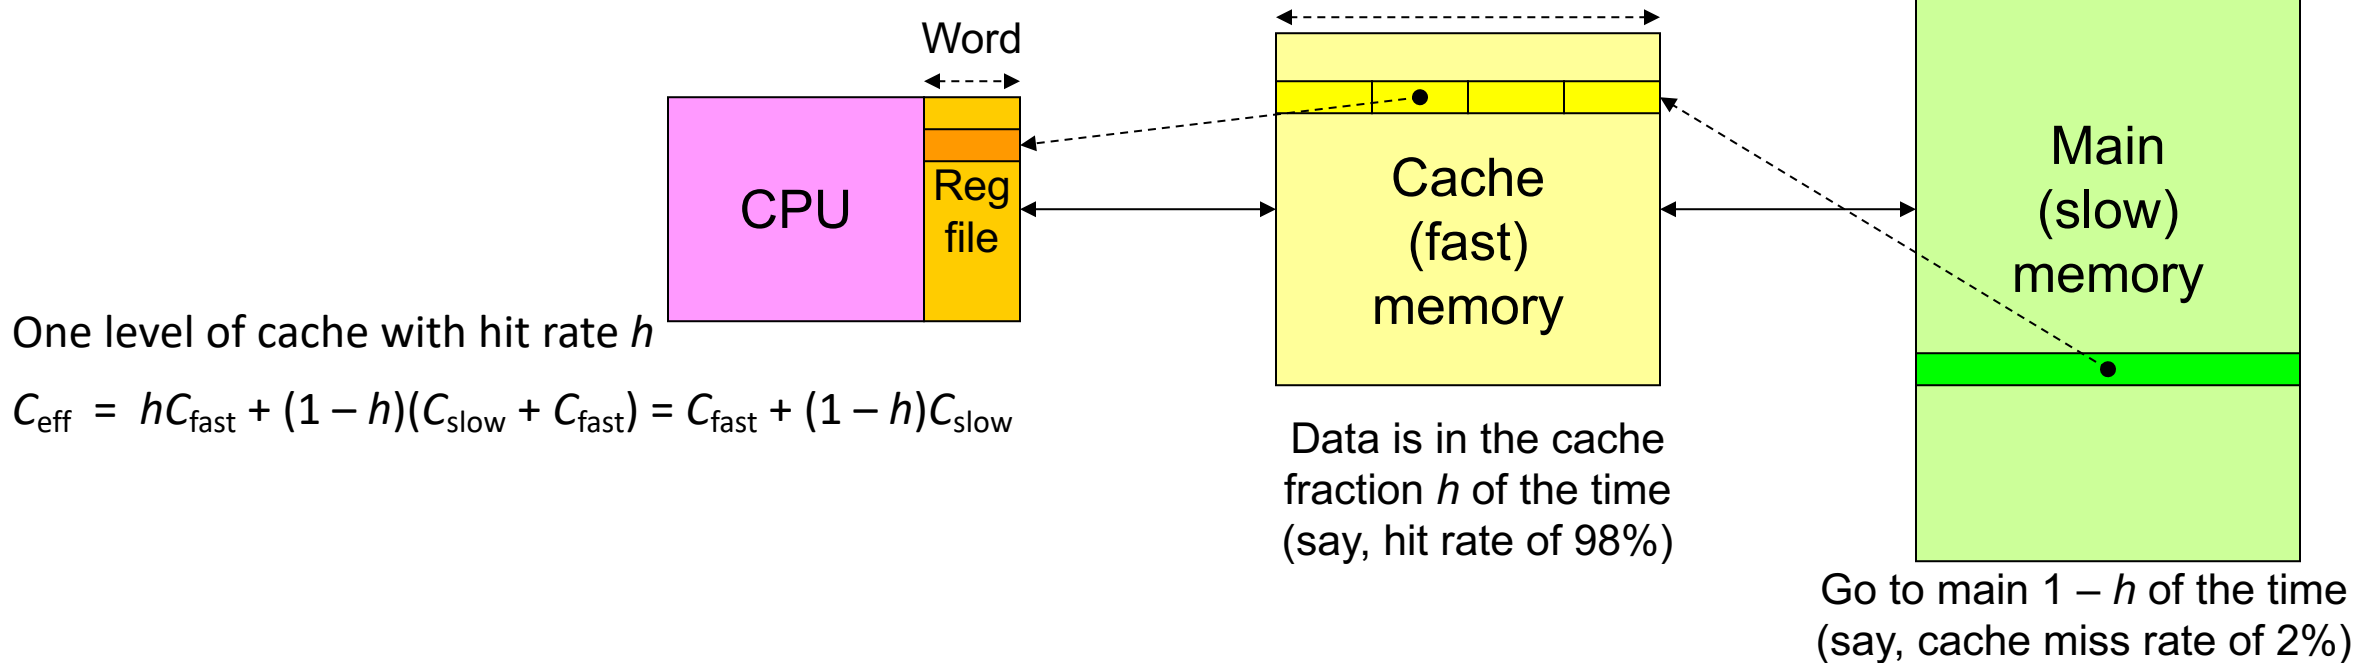
- The widening speed gap between CPU and main memory
 - Processor operations take of the order of 1 ns
 - Memory access requires 10s or even 100s of ns
- Memory bandwidth limits the instruction execution rate
 - Each instruction executed involves at least one memory access
 - Hence, a few to 100s of MIPS is the best that can be achieved
 - A fast buffer memory can help bridge the CPU-memory gap
 - The fastest memories are expensive and thus not very large
 - A second (third?) intermediate cache level is thus often used



Cache Memory –The Operative Terms

- Hit and Miss

Cache is transparent to user;
transfers occur automatically

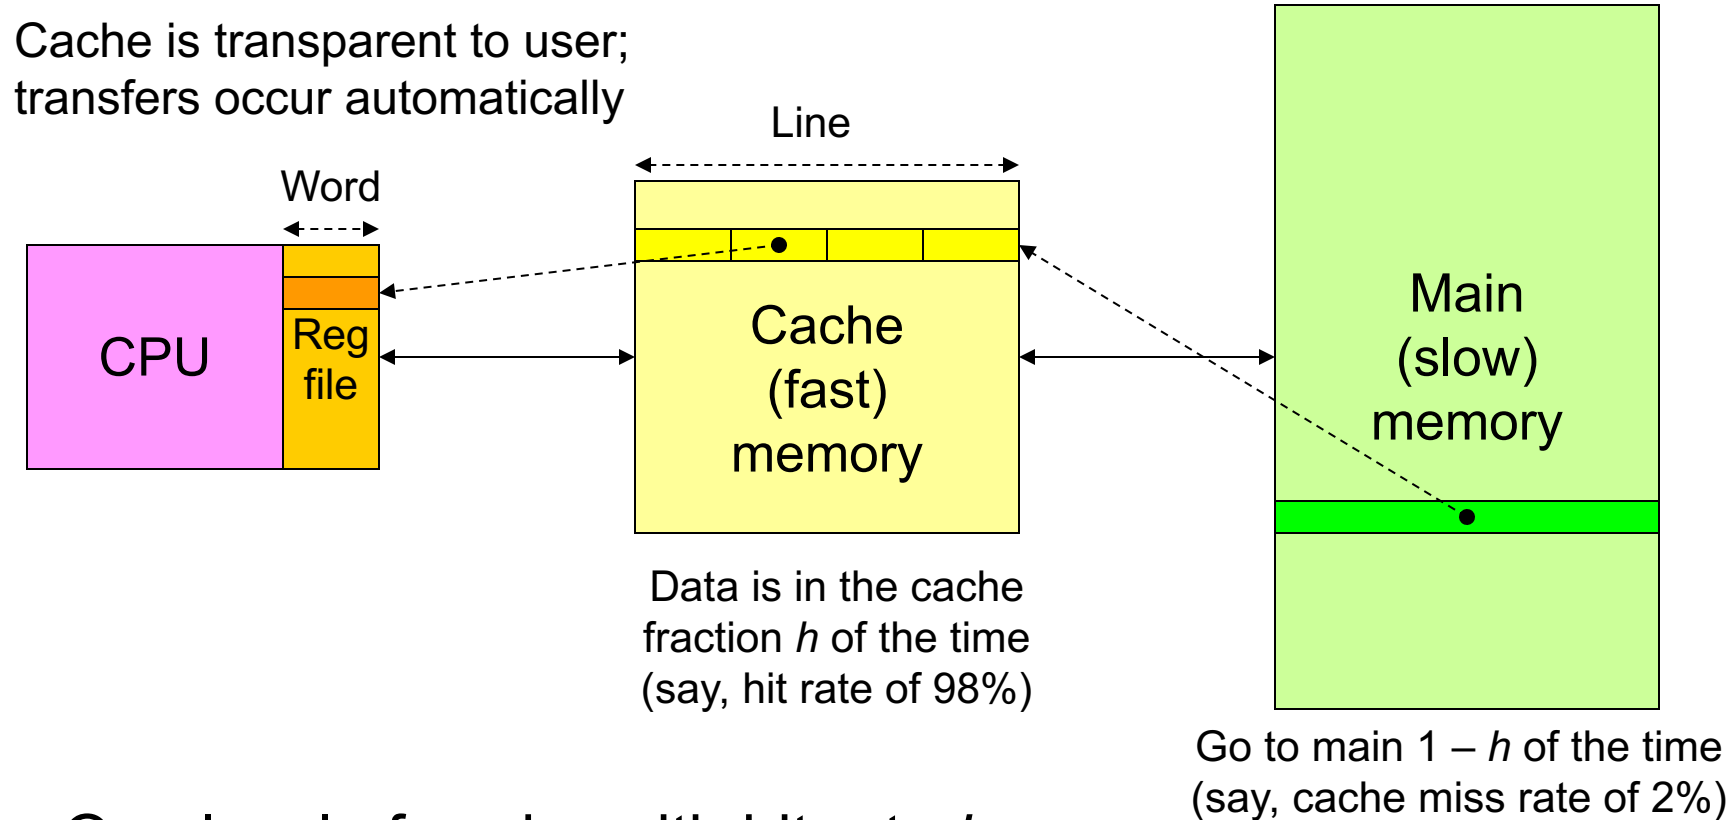


Compulsory misses: With *on-demand fetching*, first access to any item is a miss. Some “compulsory” misses can be avoided by prefetching.

Capacity misses: We have to oust some items to make room for others. This leads to misses that are not incurred with an infinitely large cache.

Conflict misses: Occasionally, there is free room, or space occupied by useless data, but the mapping/placement scheme forces us to displace useful items to bring in other items. This may lead to misses in future

Cache, Hit/Miss Rate, and Effective Access Time



One level of cache with hit rate h

$$C_{\text{eff}} = hC_{\text{fast}} + (1 - h)(C_{\text{slow}} + C_{\text{fast}}) = C_{\text{fast}} + (1 - h)C_{\text{slow}}$$

Cache Memory Design Parameters

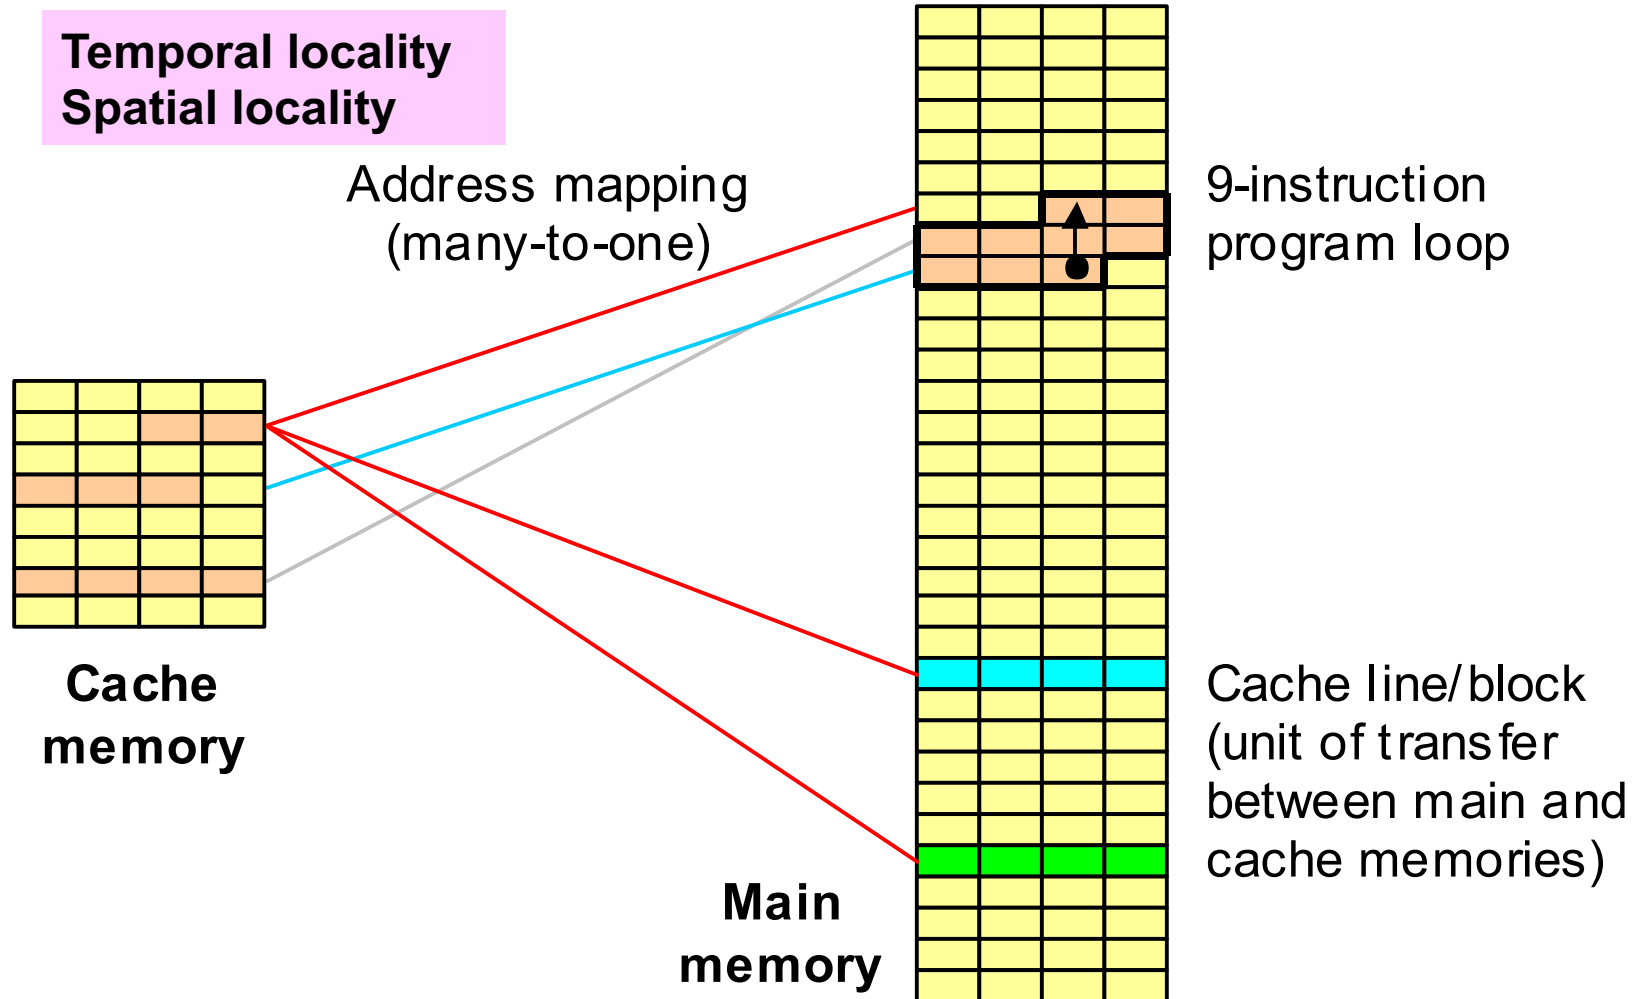
- *Cache size*
 - In bytes or words
 - A larger cache can hold more of the program's useful data but is more costly and likely to be slower.
- *Block or cache-line size*
 - Unit of data transfer between cache and main
 - With a larger cache line, more data is brought in cache with each miss.
 - This can improve the hit rate but also may bring low-utility data in.
- *Placement policy.*
 - Determining where an incoming cache line is stored
 - More flexible policies imply higher hardware cost and may or may not have performance benefits (due to more complex data location).
- *Replacement policy.*
 - Which of several existing cache blocks should be overwritten.
 - Typical policies: choosing a random or the least recently used block.
- *Write policy.*
 - *write-through*
 - Updates to cache words are immediately forwarded to main
 - *write-back or copy-back*
 - modified blocks are copied back to main if and when they must be replaced

Locality of Reference

- Why does cache work?

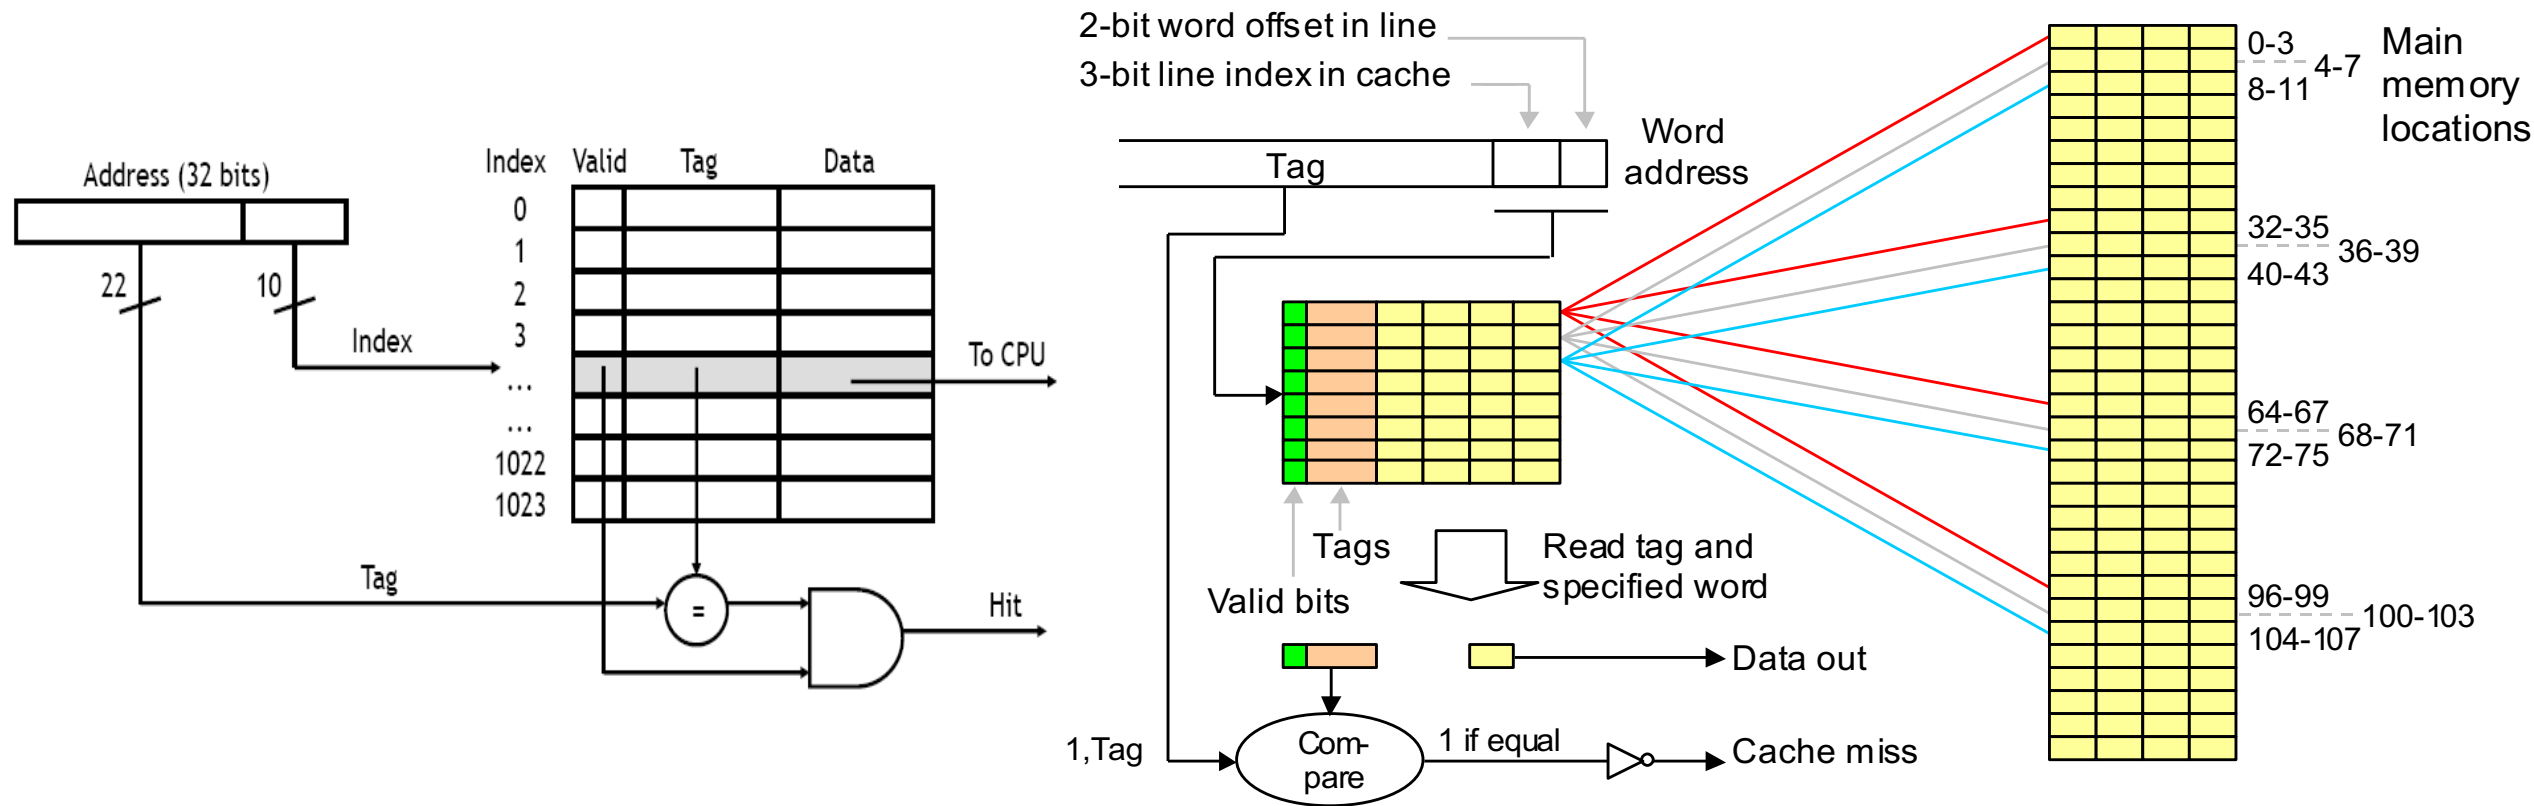
```
sum= 0;
for (i= 0; i< MAX; i++)
    sum= sum+ f(i);
```

```
sum= 0;
for (i= 0; i< MAX; i++)
    sum= sum+ a(i);
```



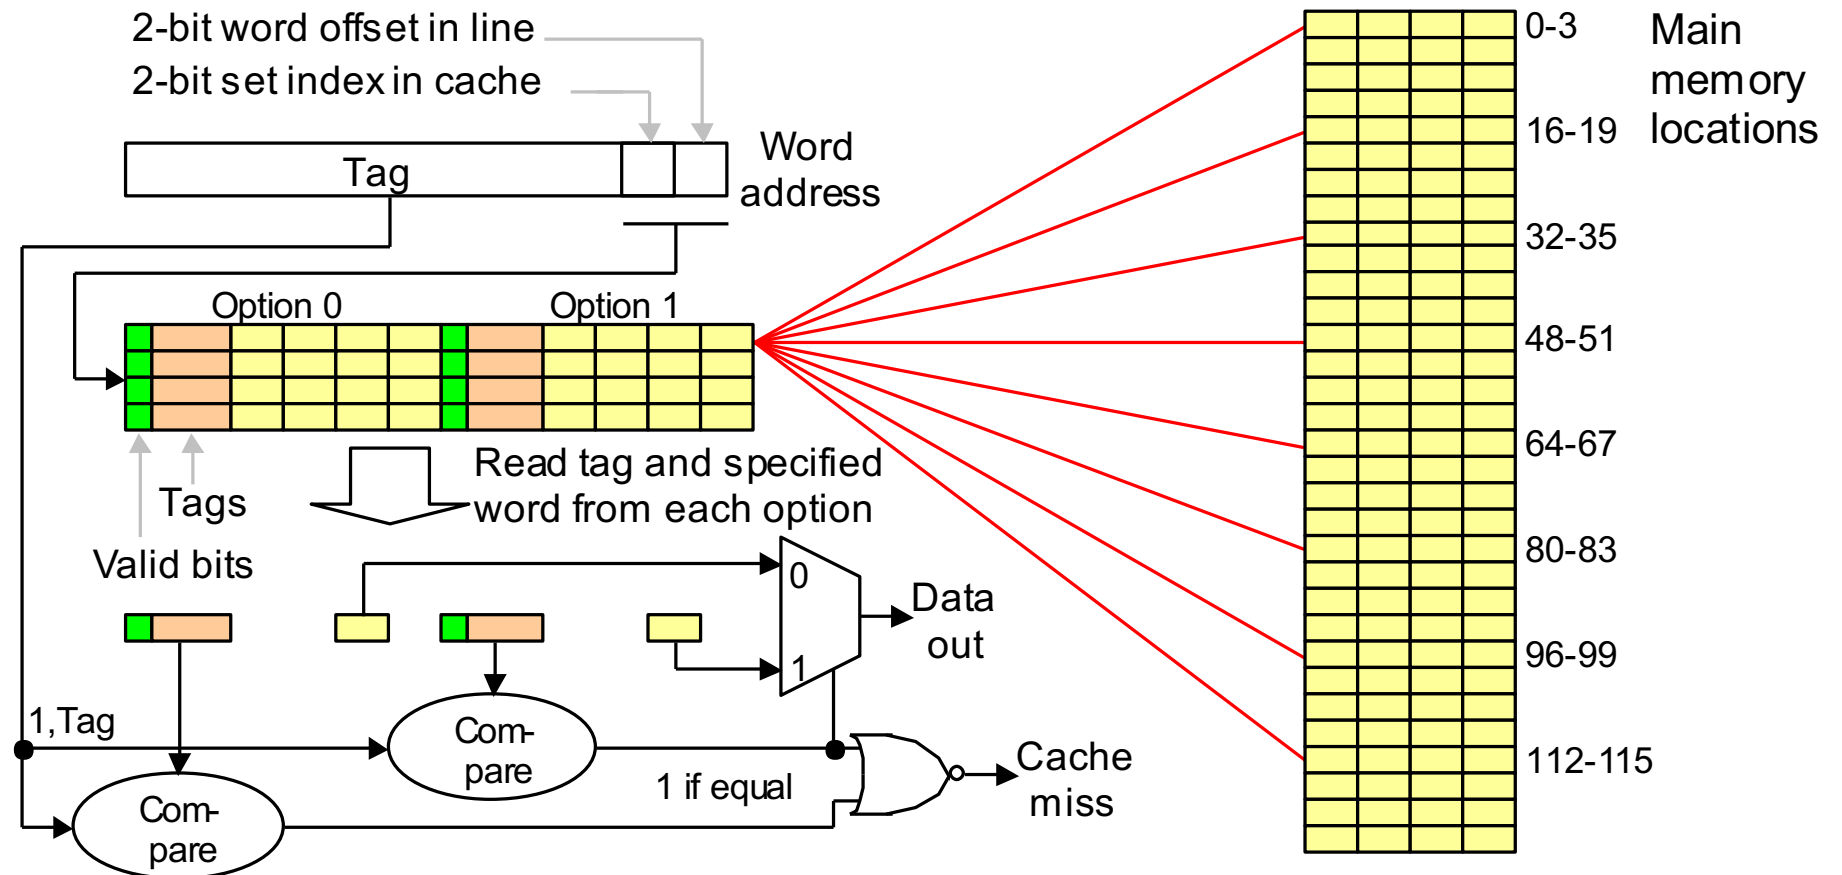
Direct Mapped Cache

- Caches are divided into blocks, which may be of various sizes
 - Each main memory address maps to exactly one cache block
 - $i \bmod 2^k$



Set Associative Cache

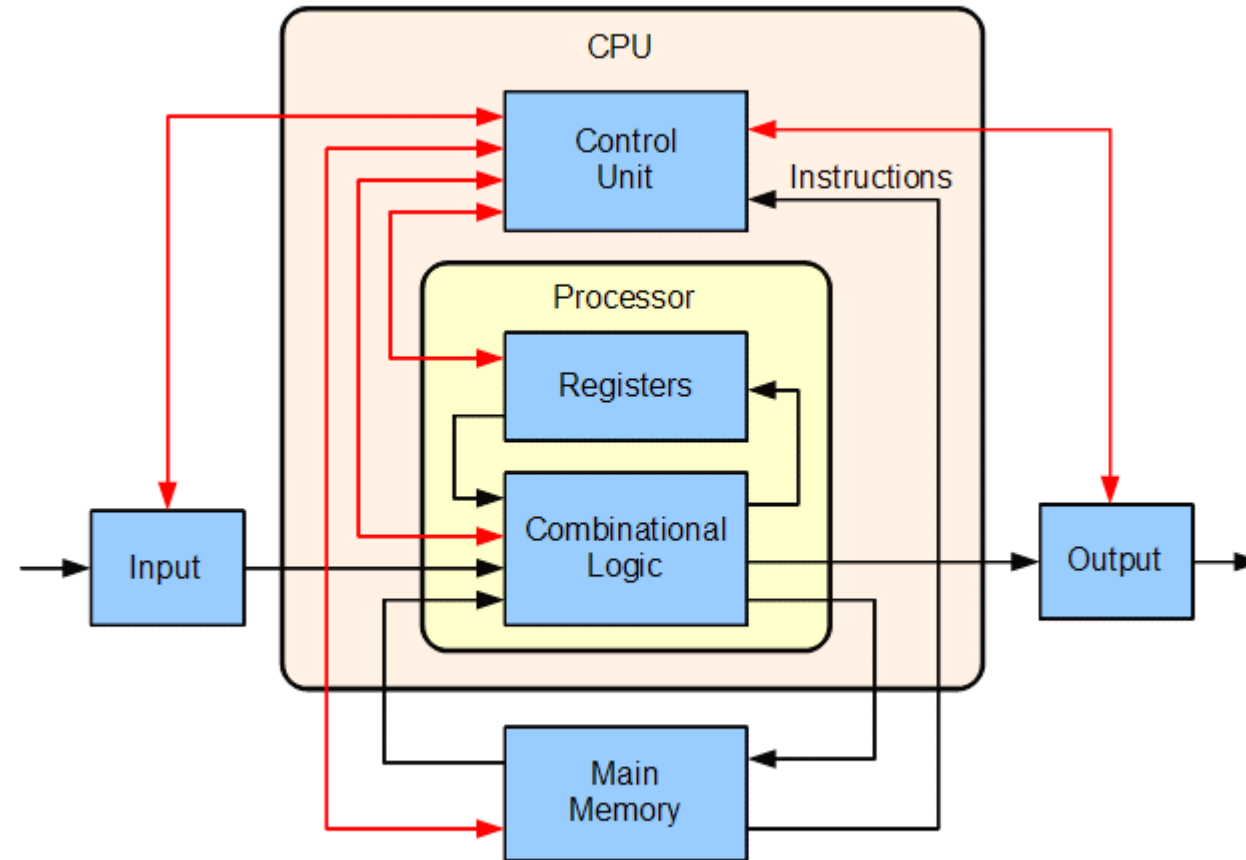
- Fully associative
 - Data to be stored in any unused cache block
- Two-way set-associative cache
 - Holds 32 words of data within 4-word lines and 2-line sets.



The Write Problem

- Architecture
 - Split cache: separate instruction and data caches (L1)
 - Unified cache: holds instructions and data (L1, L2, L3)
- What happens on a write
 - Write-through slows down the cache to allow main to catch up
 - Write-back or copy-back is less problematic, but still hurts performance due to two main memory accesses in some cases
 - Provide write buffers for the cache so that it does not have to wait for main memory to catch up

Simulator –gem5



Introduction

- Gem5 is the fusion of two projects
 - GEMS
 - detailed and flexible memory system model
 - Includes support for multiple cache coherence protocols and interconnect models
 - developed @ The University of Wisconsin Madison
 - M5
 - Highly configurable simulation framework to support multiple ISAs, and diverse CPU models
 - developed @ The University of Michigan
- System simulator
 - Good support of complex components interactions (OS / CPU / Caches / Devices / ...)
 - Accuracy depends on the model completeness
 - Lot of components available out-of-the-box (CPUs, memories, I/Os, ...)

gem5

- Good for

- Architectural exploration

- Gem5 provides a fast and easy framework to interconnect hardware components and evaluate them !
 - Hardware/software performance evaluation ?
 - Gem5 has a good support of various ISA and allows for realistic HW/SW performance evaluation

- Not Good for


- Hardware/software **verification** ?
 - RTL functional verification is much more accurate !
 - Software **development** and verification ?
 - Faster technologies are available through binary-translation (e.g. QEMU, OVP)

Simulation modes

- Full-system (FS)
 - Models **bare-metal hardware**
 - Includes the various specified devices, caches, ...
 - Boots an **entire OS** from scratch
 - Gem5 can boot Linux (several variants) or Android out of-the-box
- Syscall Emulation (SE)
 - Runs a **single** static **application**
 - System calls **are emulated or forwarded** to the host OS
 - Lot of **simplifications** (address translation, scheduling, no pthread ...)

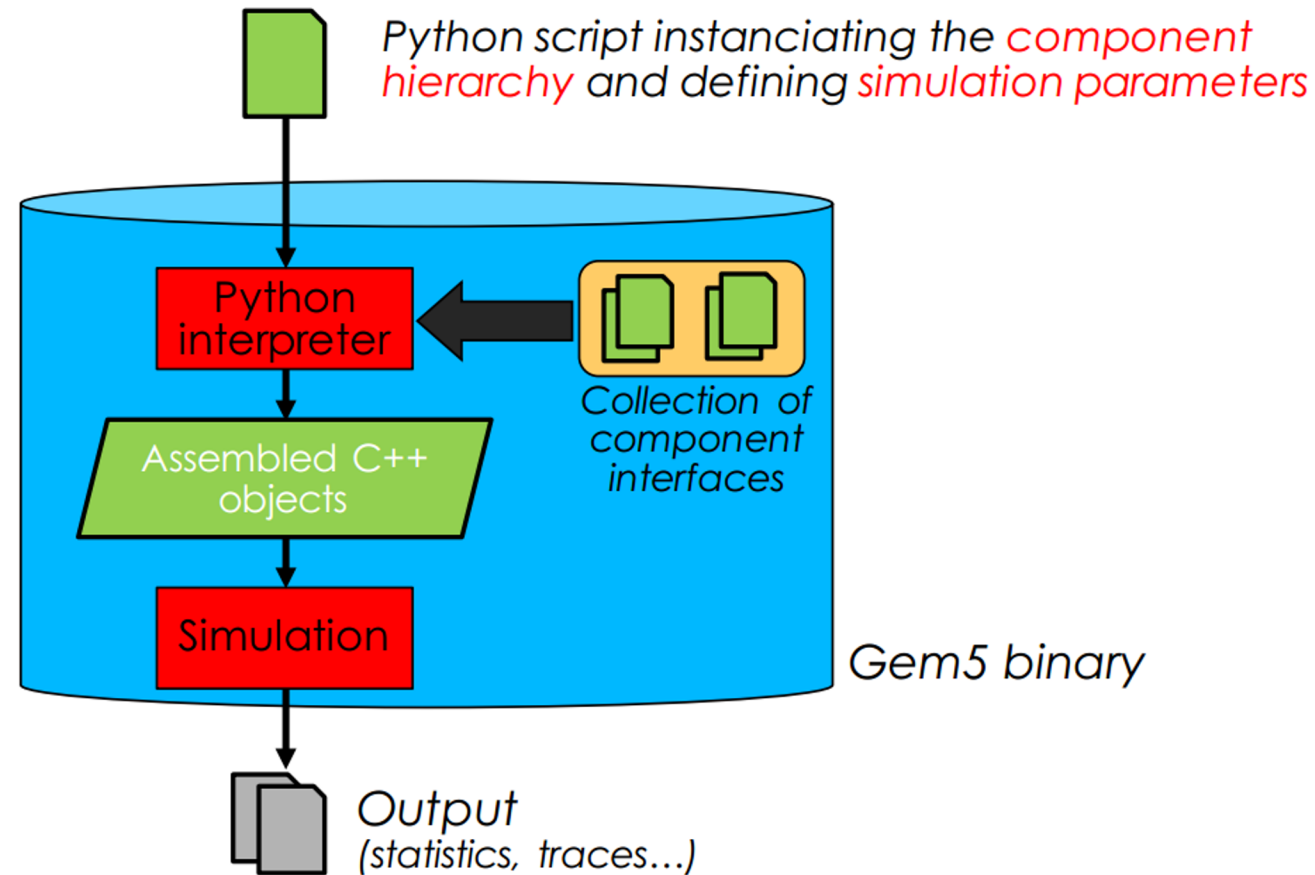
Behind the scene of a simulation

Compilation of the simulator



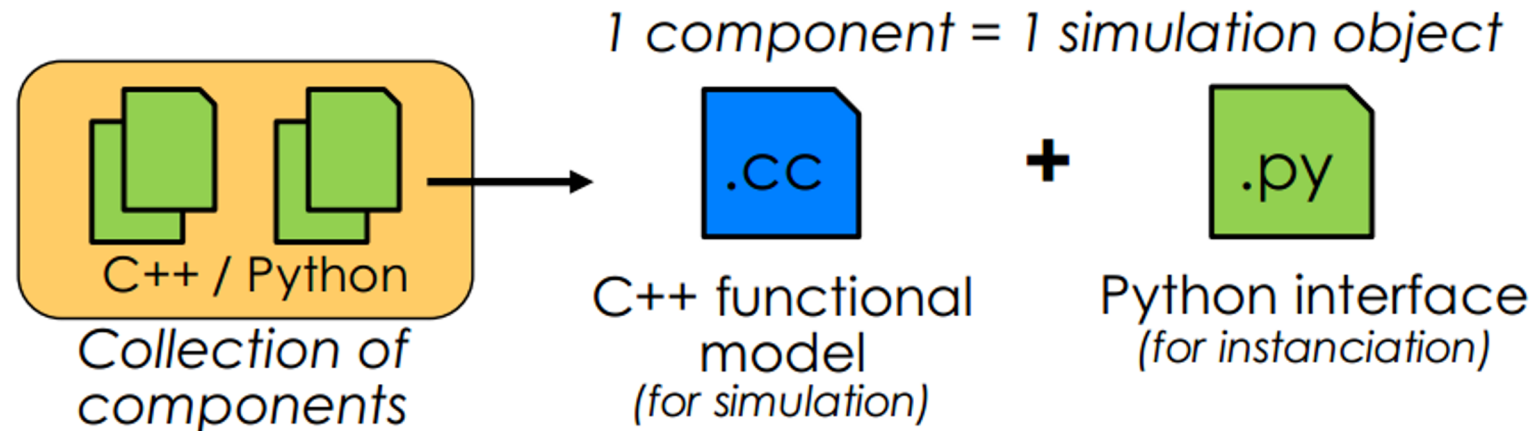
Behind the scene of a simulation

Compilation of the simulator




What's under the hood ?

Simulation objects



- SimObjects follow a strict C++ class hierarchy for easier extension with code reuse



What's under the hood ?

Events

- Gem5 is event-driven
 - Discrete event timing model
- Simulation objects schedule events for the next cycle after a specific time elapsed

What's under the hood ?


CPU Models

- **Supports**: Alpha, ARM, MIPS, Power, SPARC, and x86
- Configurable CPU models : Supports 3 CPU models
 - Simple Atomic/Timing
 - **Fast** CPU model
 - InOrder
 - Detailed **pipelined in-order** CPU model
 - O3
 - Detailed **pipelined out-of-order** CPU model
- Supports a domain specific **language** to **represent ISA** details

What's under the hood ?

Memory System

- Models a system running heterogeneous applications
 - running on **heterogeneous processing** tiles
 - using **heterogeneous memories** and interconnect



What's under the hood ?


Memory System

- Two memory systems in Gem5
 - Classic
 - All components **instantiated** in a **hierarchy** along with CPUs, etc.
 - **Only MOESI** coherence protocol
 - Fast, but less detailed than Ruby
 - Ruby
 - Detailed simulation model of **various cache hierarchies**
 - **Various** cache coherence **protocols** (MESI, MOESI, ...)
 - Interconnection networks

What's under the hood ?

Memory ports

- Memory ports are present on every MemObject
 - They model **physical memory connections**
 - You interconnect them during the hierarchy instantiation
 - E.g. a CPU data bus to a L1 cache
 - **1 master** port always connects to **1 slave** port
- Data is exchanged atomically as **packets**



What's under the hood ?

Memory ports

- 3 types of transport interfaces for packets
 - Functional
 - **Instantaneous** in a single function call
 - Caches and memories are updated **at once**
 - Atomic
 - **Instantaneous**
 - **Approximate latency**, but no contention nor delay
 - Timing
 - Transaction split into **multiple phases**
 - **Models all timing** in the memory system

Nice feature

- Checkpointing
 - **Snapshot** the relevant system state and **restore** it **later**
- Fast-forward
 - Idea is to start the simulation in **atomic mode** and switch over to **detailed mode** for **important** simulation period